

Crane Cluster Manual

University of the Witwatersrand
2017

Written by

Or Hanoach

Joshua Bruton

Meir Rosendorff

Nathan Michlo (Formatting)

Group Instructor

Atif Muhammad

| | |
|---|-----------|
| ABOUT THE CRANE CLUSTER..... | 3 |
| HEAD NODE..... | 3 |
| NODES..... | 3 |
| NFS - NETWORK FILE SYSTEM | 4 |
| INSTALLATION | 4 |
| ENABLE AND START | 4 |
| SERVER CONFIGURATION | 4 |
| CLIENT CONFIGURATION | 5 |
| MOUNTING | 5 |
| FIREWALL | 5 |
| MPI USING NFS (V1) | 6 |
| INSTALLATION | 6 |
| ADDING PATHS | 6 |
| CONFIGURATION AND MODULES | 7 |
| TESTING | 7 |
| MPI USING NFS (V2) | 8 |
| DOWNLOAD..... | 8 |
| INSTALLATION | 8 |
| ADDING PATHS | 9 |
| CONFIGURATION AND MODULES | 9 |
| TESTING | 10 |
| DNS SERVER (DOMAIN NAMING SYSTEM) | 11 |
| INSTALLATION | 11 |
| CONFIGURATION | 11 |
| FORWARD ZONE..... | 12 |
| REVERSING ZONE..... | 13 |
| ENABLING..... | 13 |
| FIREWALL | 13 |
| PERMISSIONS | 14 |
| CHECKING | 14 |
| TESTING | 15 |
| QUOTA..... | 16 |
| PRE-SETUP | 16 |
| SETTING A QUOTA | 16 |
| TORQUE | 18 |
| HEAD NODE INSTALLATION..... | 18 |
| NODES INSTALLATION..... | 20 |
| INSTALLATION ON CLIENTS..... | 22 |
| ENABLE TORQUE AS A SERVICE | 22 |
| COMMANDS | 23 |
| TESTING | 25 |
| LDAP (LIGHTWEIGHT DIRECTORY ACCESS PROTOCOL) | 27 |
| INSTALLATION ^[1] | 27 |
| USER ACCOUNT MIGRATION..... | 32 |
| FIREWALL CONFIGURATION | 33 |

| | |
|--|-----------|
| LDAP CLIENT CONFIGURATION ^[2] | 33 |
| NFS SERVER CONFIGURATION | 35 |
| AUTOMOUNTER CLIENT CONFIGURATION | 35 |
| CREATE NEW USER | 35 |
| CREATE NEW GROUP | 35 |
| WAKE ON LAN | 37 |
| CONFIGURATION | 37 |
| EXECUTION | 37 |
| TESTING | 37 |
| FORMAT AND MOUNTING HDD..... | 38 |
| FORMAT..... | 38 |
| REMOVE PARTITIONS..... | 38 |
| MOUNT | 38 |
| ADDITIONAL SCRIPTS | 40 |
| NETWORK_CONFIG..... | 40 |
| PERFORM | 40 |
| PING_ALL | 40 |
| WITS_PROXY | 40 |
| WAKE_CRANES | 41 |
| LDAP SCRIPTS | 41 |
| GANGLIA ^[1] ^[2] | 42 |
| INSTALLATION | 42 |
| CONFIGURATION | 42 |
| CLONING NODES | 44 |
| CREATE BOOTABLE FLASH DRIVE | 44 |
| CREATING IMAGE ^[1] | 44 |
| RESTORING AN IMAGE ^[4] | 44 |
| NEW STUFF | 46 |
| SET HOSTNAME | 46 |
| NETWORKING | 46 |
| CREATE SSH-KEY | 46 |
| CHANGE DATE ON MACHINES | 47 |
| ON ADDITION OF NODES: | 47 |
| REFERENCES | 48 |

About the Crane Cluster

The Crane Cluster is a computer cluster made up of one head node and many separate nodes it uses. As of the writing of this manual, the Crane cluster consists of the head node and 19 compute nodes, located in MSL building, second floor, research lab 2.

Head Node

The head node has 1 physical Ethernet card split logically into two cards - one for the internal cluster LAN and one for internet connectivity through the wits network. The external internet connection goes through the Wits proxy and as such gets its IP using the Wits DHCP. The internal LAN connection uses a static IP.

| | |
|------------------------------|--------------|
| Hostname | crane |
| LAN IP | 10.0.0.101 |
| Internet IP (as of writing): | 10.10.187.31 |

Compute Nodes

| | | |
|-----------------------------|----------|---|
| Hostname | crane# | (Where # is the number of the node) |
| LAN IP | 10.0.0.# | (Where # is the number of the node) |
| Internet IP (as of writing) | n/a | (Nodes only have a static IP for the internal LAN, they do not have access to the internet directly.) |

NFS - Network File System

We use an NFS server to have a shared folder for head node and all the other nodes.

The NFS is setup with the folder on the server being `/server/folder/` and the folder on the clients being `/client/folder/`

In our case the NFS folder is located on head node as the folder `/craneNest`. Physically this folder is located on a separate 1TB HDD mounted to head node from sdb (see Format and Mounting HDD).

Installation

Run the following command on both the server and the clients:

```
$ yum install nfs-utils
```

Enable and Start

Run the following commands on both the server and the clients:

```
$ systemctl enable rpcbind
$ systemctl enable nfs-server
$ systemctl enable nfs-lock
$ systemctl enable nfs-idmap
$ systemctl start rpcbind
$ systemctl start nfs-server
$ systemctl start nfs-lock
$ systemctl start nfs-idmap
```

Server Configuration

Create the public folder that will be shared and give it full permissions:

```
$ mkdir /server/folder
$ chmod 777 /server/folder
```

Add the following to `/etc/exports` to define who has permissions to access the public folder and what their permissions are. (<client IP> must be replaced by the range of client IPs, the part in the brackets defines the types of permissions. NB the only space is between the folder path and the client IP.)

```
/server/folder/ <client IP>(rw,sync,no_root_squash,no_all_squash)
```

Restart NFS with:

```
systemctl restart nfs-server
```

Client Configuration

Create the folder in which the data inside `/server/folder` will be displayed:

```
$ mkdir /client/folder
```

Mounting

Use the following to mount the `/server/folder` inside the `/client/folder` replacing `<server IP>` with the server's IP.

```
$ mount -t nfs <server IP>:/server/folder /client/folder
```

NOTE: If the previous command times-out, see [firewall section](#) and then return to this section and try again.

In order to set it to mount on boot add the following to the `/etc/fstab` file replacing `<server IP>` with the server's IP:

```
<server IP>:/server/folder/ /client/folder/ nfs rw,sync,hard,intr 0 0
```

In order to confirm correct mount enter the `df` command and you should see your folder on the list.

Firewall

Uncomment the following in `/etc/sysconfig/nfs` order to set default ports

```
MOUNTD_PORT=port
STATD_PORT=port
LOCKD_TCP_PORT=port
LOCKD_UDP_PORT=port
```

Use the following command for all the ports you just uncommented in `/etc/sysconfig/nfs` as well as ports 111 and 2049 in order to allow nfs through the firewall where `<port>` is the port number)

```
$ firewall-cmd --permanent --add-port=<port>/tcp
```

After allowing all the ports through the firewall use the following command to restart the firewall:

```
$ firewall-cmd -reload
```

MPI using NFS (v1)

MPI is a protocol used to pass messages, over a network, between programs running in parallel across a host or cluster.

Open MPI is an MPI implementation used to handle and run these such jobs and programs.

Below, it will be installed and setup in the shared NFS folder `/craneNest/` in the `apps/` directory.

Installation

Create the folder for Open MPI in apps using:

```
$ mkdir /craneNest/apps/openmpi
```

Go to the download location of Open MPI and run the following:

```
$ ./configure --prefix = /craneNest/apps/openmpi
$ make
$ make install
```

Adding PATHs

You need to add `/craneNest/apps/openmpi/bin` and `/craneNest/apps/openmpi/lib` to the `$PATH` of all the nodes on the server.

You can do this without a restart by typing the following commands:

```
$ export PATH=$PATH:/craneNest/apps/openmpi/bin
$ export PATH=$PATH:/craneNest/apps/openmpi/lib
```

However, to make the change permanent you must add the above commands to `/home/<user_name>/.bashrc` and `/etc/profile`. (in our case the users are crane and root)

Adding to profile will make it activate when a user logs-in and adding it to: `.bashrc` will run it when the computer boots.

Configuration and Modules

Enter the command

```
$ which mpirun
```

If it does not return a path perform the following or run the following commands:

```
$ module avail
```

```
$ add module <path>
```

 where <path> is the path seen at the end of the previous output.

Enter the command `$ which mpirun` to confirm it outputs a path.

Enter the command `$ mpiexec --version` and it should return the details of your Open MPI version if the paths are configured correctly.

Testing

To test Open MPI, a host file must first be created on which is simply a list of IP addresses or domain names that you want to run Open MPI on.

Run Open MPI with the following command:

```
$ mpirun -hostfile <name of hostfile> -np 1 echo Hello
```

If Open MPI was configured correctly this should output Hello in the terminal of the computers specified in the `hostfile`.

Test this on several the nodes to confirm the `$PATH`'s are configured correctly everywhere.

To test if the processes are being divided amongst the nodes correctly run the following:

(The sleep command will put the running user to sleep for the number specified in seconds. The `-np` flags specifies the number of processors to be used)

```
$ mpirun -hostfile <name of hostfile> -np 1 echo sleep 30
```

Before the process is completed enter the following commands into any of the nodes specified in the host file:

```
$ top | grep <name of user>
```

This should produce at least one sleep process (depending on the number of processors specified and one orted process (orted is the method MPI uses to run.)

NOTE

If this works then Open MPI is running correctly.

MPI using NFS (v2)

MPI is a protocol used to pass messages, over a network, between programs running in parallel across a host or cluster.

Open MPI is an MPI implementation used to handle and run these such jobs and programs.

Below, it will be installed and setup in the shared NFS folder `/craneNest/` in the `apps/` directory.

Download

Download onto a USB stick the latest release version of Open MPI and copy it to the head node.

Installation

Create the folder for Open MPI in `/apps` using:

```
$ mkdir /craneNest/apps/openmpi
```

Go to the download location of Open MPI and run the following:

```
$ gunzip -c openmpi-2.0.2.tar.gz | tar xf -  
$ cd openmpi-2.0.2  
$ ./configure --prefix=/usr/local  
$ make all install
```

You might need to update automake if `make all install` does not work:

```
$ wget http://ftp.gnu.org/gnu/automake/automake-1.15.tar.gz  
$ tar xvzf automake-1.15.tar.gz  
$ cd automake-1.14  
$ ./configure  
$ make  
$ sudo make install
```

Adding PATHs

You need to add `craneNest/apps/openmpi/bin` and `craneNest/apps/openmpi/lib` to the `$PATH` of all the nodes on the server.

You can do this without a restart by typing the following commands:

```
$ export PATH=$PATH:/craneNest/apps/openmpi/bin
$ export PATH=$PATH:/craneNest/apps/openmpi/lib
```

However, to make the change permanent you must add the above commands to `/home/<user_name>/.bashrc`. (in our case the users are `crane` and `root`)

Adding to profile will make it activate when a user logs-in and adding it to: `.bashrc` will run it when the computer boots.

Configuration and Modules

Enter the command

```
$ which mpirun
```

If it does not return a path perform the following:

```
$ module avail
```

`$ add module <path>` where `<path>` is the path seen at the end of the previous output.

Enter the command `which mpirun` to confirm it outputs a path.

Enter the command `$ mpiexec --version` and it should return the details of your Open MPI version if the paths are configured correctly.

Testing

(If you run into problems with the firewall you might need to turn it off until we figure out how to make exceptions that work “systemctl stop firewalld”, “systemctl disable firewalld”)

To test Open MPI, a host file must first be created on which is simply a list of IP addresses or domain names that you want to run MPI on.

Run MPI with the following command:

```
$ mpirun -hostfile <nameOfHostfile> -np 1 echo Hello
```

If Open MPI was configured correctly this should output “Hello” in the terminal of the computers specified in the `hostfile`.

Test this on several nodes to confirm the `$PATH`'s are configured correctly everywhere.

To test if the processes are being divided amongst the nodes correctly run the following:

(The `hostname` command will output the running nodes hostname. The `-np` flags specifies the number of processors to be used)

```
$ mpirun -hostfile nameOfHostfile -np 1 sleep hostname
```

NOTE

If this works then Open MPI is running correctly.

DNS Server (Domain Naming System)

A DNS Server is a server that pairs numbered IP addresses to named network addresses, and creates a domain for a network to work on. We used the head node as our DNS server.

Installation

Install bind by running:

```
$ yum install bind bind-utils -y
```

Configuration

Open `/etc/named.conf` and once there comment out `listen-on-v6 port 53 { ::1; };` and add your DNS server IP to `listen-on port 53 {<ip>;};`

In `allow-query` add the IP range of your slaves. e.g. `localhost; 10.0.0.0/24;` The `allow-transfer` command can be used to configure a secondary backup DNS, if you don't have one, comment it out.

At the bottom of your file, under the already existing zone, instantiate your forward and reverse DNS lookup zones. For example, using `zonename` as the zone's name.

```
zone "zonename" IN {
type master;
file "forward.zonename";
allow-update { none; };
};

zone "x.x.x.in-addr.arpa" IN {                                // 'x.x.x' is the first
3 octets of your DNS IP in reverse e.g. '0.0.10'.
type master;
file "reverse.zonename";
allow-update { none; };
};
```

Now you must configure and create the aforementioned zones.

Save and exit `named.conf`

Open `/var/named/` and create `forward.zonename` and `reverse.zonename` in the folder.

Forward Zone

Add the following to the `forward.zonename` file using the number of nodes required and replacing `nameofDNSHost`, `root_zonename`, `IPofDNS` and `IPofNode` with the relevant information.

NOTE: Make sure to include all your nodes in this file.

```
$TTL 86400
@   IN  SOA  nameofDNSHost.  root.zonename. (
        2011071001  ;Serial
        3600        ;Refresh
        1800        ;Retry
        604800      ;Expire
        86400       ;Minimum TTL
)
@   IN  NS   nameofDNSHost.
@   IN  A    IPofDNS
@   IN  A    IPofNode1
@   IN  A    IPofNode2
nameofDNSHost  IN  A          IPofDNS
nameOfNode1   IN  A          IPofNode1
nameOfNode2   IN  A          IPofNode2
```

Reversing Zone

Add the following lines to the `reverse.zonename` file using the number of nodes required and replacing `nameofDNS`, `root_zonename`, `IPofDNS` and `IPofNode` with the relevant information. NB make sure to include all your nodes in this file.

```
$TTL 86400
@ IN SOA      nameofDNShost. root.zonename. (
    2011071001 ;Serial
    3600        ;Refresh
    1800        ;Retry
    604800      ;Expire
    86400       ;Minimum TTL
)
@ IN NS       nameofDNShost.
@ IN PTR      zonename.
nameofDNShost IN A    iPoFDNS
nameofNode1   IN A    iPoFnode1
nameofNode2   IN A    iPoFnode2
lastoctetofDNSIP IN PTR      nameofDNShost
lastoctetofNode1IP IN PTR      nameofNode1
lastoctetofNode2IP IN PTR      nameofNode2
```

Enabling

You can now turn on the DNS by running the following

```
$ systemctl enable named          (Will start named on boot)
$ systemctl start named           (Will start named now)
```

Firewall

Configure the firewall to allow DNS on port 53 with the following:

```
$ firewall-cmd --permanent --add-port=53/tcp
$ firewall-cmd --permanent --add-port=53/udp
```

Reload the firewall for the new settings to take effect with:

```
$ firewall-cmd --reload
```

Permissions

Permissions, ownership and SELinux configuration is done with the following:

```
$ chgrp named -R /var/named
$ chown -v root:named /etc/named.conf
$ restorecon -rv /var/named
$ restorecon /etc/named.conf
```

Checking

Run the following if there is no output continue.

```
$ named-checkconf /etc/named.conf
```

Zones checks replace zonename with your zone's name, if it returns OK, they are okay.

```
$ named-checkzone zonename /var/named/forward.zonename
$ named-checkzone zonename /var/named/reverse.zonename
```

Add the DNS IP to your `ifcfg` file so it knows who to contact. (if you do not use `eno1`, replace it with your internet interface.)

```
vim /etc/sysconfig/network-scripts/ifcfg-eno1
```

~~Add DNS="IPofDNS" to that ifcfg-eno1 and then save and exit.~~

Go to `/etc/resolv.conf` and add at the beginning:

```
search zonename
nameserver IPofDNS
```

Comment out (#) all the existing lines except for the first two nameservers (other than your DNS).

note: You must add this to all the `resolv.conf` files of the nodes as well, BUT you can only search one zone per node.

Now restart the network and DNS.

```
$ systemctl restart network
$ reboot
```

Check your `resolv.conf` file and make sure your changes are still there. If not, add them again and use the command `$ chatter +i /etc/resolv.conf` to make the file non-editable. To undo this, use `$ chatter -i /etc/resolv.conf`.

Testing

Use: `nslookup zonename` and you should receive the following:

```
Server:          IPofDNS
Address:         IPofDNS#53

Name: nameofZone
Address: IPofDNS
Name: nameofZone
Address: IPofNode1
Name: nameofZone
Address: IPofNode2
.
.
.
```

Reboot all nodes

Now ping an active node using its name (not its IP-address) and you should get a response.

NOTE

If both tests work the DNS is running.

Quota

Quota allows you to limit the size of folders for certain users

Pre-Setup

To setup perform the following:

add “usrquota,grpquota” to the folder you want to add a quota to in /etc/fstab

For example, in our case we add a quota to craneNest by editing /etc/fstab in the following way:

```
/craneNest/      home  ext3  defaults,usrquota,grpquota 1 2
```

Remount the folder to apply the settings:

```
$ mount -o remount the folder / reboot
```

Use the following to confirm your folder has quota enabled

```
$ mount | grep quota
```

Now switch quota on for that folder

```
$ quotaon /craneNest
```

Setting a Quota

Now we set the actual quota for the folder

Use the following command to set it for the user replacing USER with the user’s name

```
$ edquota -u USER
```

and to set it for groups, replacing GROUP with the group’s name we use

```
$ eqquota -g GROUP
```

This will open a window that looks like the following for a user named Jack

Disk Quotas for user Jack (uid 1001):

| Filesystem | blocks | soft | hard | inodes | soft | hard |
|-------------------------|--------|------|------|--------|------|------|
| /dev/mapper/centos-home | 0 | 5500 | 6000 | 0 | 0 | 0 |

By editing the first two values underneath hard and soft we can change the quota on the disk. When the user's usage reaches the soft value, he will receive a warning and when he reaches the hard value it will no longer allow him to create new files.

The `inodes` field allows us to limit the number of files the user can make.

`$ edquota -t` allows you to configure grace period for soft limit after which it becomes hard storage

In order to track the usage, we use the command `$ repquota -as` which will give a summary of all the different user's usages and their limits

Automation of the process

This can be done for user creation.

We can automate the process by creating a template user and copying those settings to other users upon their creation. In this case, we use the following command where `<user>` is the user we are creating and `<user prototype>` is the name of the prototype user we want to use

```
$ edquota <user> -p <user prototype>
```

Torque

Torque is a program that allows scheduling and distribution of jobs from a head node to separate nodes of a cluster.

Head Node Installation

Opening Ports

Torque requires certain ports to be open for essential communication.

- For client and pbs_mom communication to pbs_server, the default port is 15001.
- For pbs_server communication to pbs_mom, the default port is 15002.
- For pbs_mom communication to pbs_mom, the default port is 15003.

```
$ firewall-cmd --add-port=15001/tcp --permanent
$ firewall-cmd --reload
```

Verify Hostname

Make sure that the correct hostname and IP are entered correctly in /etc/hosts (in our case "crane")

To verify that the hostname resolves correctly, make sure that `hostname` and `hostname -f` report the correct name for the host.

Install required packages

```
$ yum install libtool openssl-devel libxml2-devel boost-devel gcc
gcc-c++
```

Download Torque

Go to folder you want to save Torque files at (I chose /root/programs)

Download git and copy folder from git:

```
$ yum install git
$ git clone
https://github.com/adaptivecomputing/torque.git -b 6.0.1 6.0.1
$ cd 6.0.1
$ ./autogen.sh
```

Install Torque

While in the folder 6.0.1

```
$ ./configure  
$ make  
$ make install
```

Set Torque server name

Verify that the `/var/spool/torque/server_name` file exists and contains the correct name of the server (in our case "crane").

can be done using:

```
$ echo <torque_server_hostname> > /var/spool/torque/server_name
```

Boot config

Configure the `trqauthd` daemon to start automatically at system boot

```
$ cp contrib/systemd/trqauthd.service /usr/lib/systemd/system/  
$ systemctl enable trqauthd.service  
$ echo /usr/local/lib > /etc/ld.so.conf.d/torque.conf  
$ ldconfig  
$ systemctl start trqauthd.service
```

PATHs

Make sure `/usr/local/bin` and `/usr/local/sbin` are in the `$PATH` environment variable:

Checking what is in `PATH`:

```
$ echo $PATH
```

Adding to `PATH`:

```
$ export PATH=/usr/local/bin/:/usr/local/sbin/:$PATH
```

Torque.setup Script

Initialize `serverdb` by executing the `torque.setup` script:

in 6.0.1 folder:

```
$ ./torque.setup root
```

Set nodes to be used:

Add each node to the `/var/spool/torque/server_priv/nodes` on one line in the following format (separated by lines below due to document limitations):

```
<node_host_name> np=<number_of_cores_available>  
gpus=<number_of_gpus_available_(optional)>  
<string_to_characterize_node_(optional)>
```

Configure `pbs_server` to start automatically at system boot, and then start the daemon:

```
$ qterm  
$ cp contrib/systemd/pbs_server.service /usr/lib/systemd/system/  
$ systemctl enable pbs_server.service  
$ systemctl start pbs_server.service
```

Nodes Installation

Opening Ports - done on each node:

On nodes:

```
$ firewall-cmd --add-port=15002-15003/tcp --permanent  
$ firewall-cmd --reload
```

Create packages to be installed on nodes - done on head node:

from 6.0.1 folder

```
$ make packages  
Building ./torque-package-clients-linux-x86_64.sh ...  
Building ./torque-package-mom-linux-x86_64.sh ...  
Building ./torque-package-server-linux-x86_64.sh ...  
Building ./torque-package-gui-linux-x86_64.sh ...  
Building ./torque-package-devel-linux-x86_64.sh ...
```

The package files are self-extracting packages that can be copied and executed on your production machines. Use `--help` for options.

Copy packages to nodes - done from head node:

from 6.0.1 folder

```
$ scp torque-package-mom-linux-x86_64.sh  
<mom-node>:<location-on-node>  
$ scp torque-package-clients-linux-x86_64.sh  
<mom-node>:<location-on-node>
```

Copy MOM startup script to nodes - done from head node:

```
$ scp contrib/systemd/pbs_mom.service  
<mom-node>:/usr/lib/systemd/system/
```

Install MOM packages on nodes - done on each node:

```
$ ssh root@<mom-node>
```

Go to location packages were copied to

```
$ ./torque-package-mom-linux-x86_64.sh --install  
$ ./torque-package-clients-linux-x86_64.sh --install  
$ ldconfig
```

Configure pbs_mom to start at system boot, and then start the daemon. - done on each node:

```
$ systemctl enable pbs_mom.service  
$ systemctl start pbs_mom.service
```

NOTICE

Make sure that /etc/hosts /etc/hostname /etc/sysconfig/network /var/spool/torque/server_name /var/spool/torque/server_name.new ALL have the same hostname formats in both nodes and head node!

If you change head nodes hostname you need to correct /var/spool/torque/server_name as well.

Installation on Clients

If you want Torque client commands installed on hosts other than the Torque Server Host (head node)

Copy packages to client - done from head node:

Go to 6.0.1 folder

```
$ scp torque-package-clients-linux-x86_64.sh  
<torque-client-host>:<location_client>
```

Copy the trqauthd startup script to each Torque Client Host (nodes) - done from head node:

```
$ scp contrib/systemd/trqauthd.service  
<torque-client-host>:/usr/lib/systemd/system/
```

Install packages on Clients (nodes) - done on each client (node):

```
$ ./torque-package-clients-linux-x86_64.sh --install  
$ echo /usr/local/lib > /etc/ld.so.conf.d/torque.conf  
$ ldconfig
```

Enable and start the trqauthd service - done on each client (node):

```
$ systemctl enable trqauthd.service  
$ systemctl start trqauthd.service
```

Enable Torque as a Service

To have it run in background waiting for commands

from 6.0.1 folder:

On nodes:

```
$ cp contrib/init.d/pbs_mom /etc/init.d/pbs_mom  
$ chkconfig --add pbs_mom
```

On head node:

```
$ cp contrib/init.d/pbs_server /etc/init.d/pbs_server
$ chkconfig --add pbs_server
```

Commands

| | |
|------------------------------|--|
| Create PBS Server (>= 6.0.1) | <code>\$./torque.setup <username></code> |
| Start / Stop trqauthd | <code>service trqauthd <start stop></code> |
| Start PBS Server | <code>pbs_server</code> |
| Stop PBS Server | <code>qterm</code> |
| Start MOM Server | <code>pbs_mom</code> |
| Start Scheduler | <code>pbs_sched</code> |
| Display Server Settings | <code>qmgr -c 'p s'</code> |

Give manager and operator permission to user (i.e. to be able to use qrun)

```
$ qmgr
```

Qmgr: set server managers += crane@crane

Qmgr: set server operators += crane@crane

| | |
|----------------------|--|
| Submit Job | <code>qsub <job_file> <-c enabled (optional chkpt)></code> |
| Status of Jobs | <code>qstat</code> |
| Run Job | <code>qrun <job_number></code> |
| Status of Nodes | <code>pbsnodes</code> |
| Cancel Job | <code>qdel <job_number> <-m "message" (optional)></code> |
| Place Job on Hold | <code>qhold <job_number></code> |
| Release Job Hold | <code>qrls <job_number></code> |
| Pause Running Job | <code>qsig -s STOP <job_number></code> |
| Continue Running Job | <code>qsig -s CONT <job_number></code> |
| Create Checkpoint | <code>qchkpt <job_number>vi</code> |

Job Format

Typically, a submit script is written to hold all the parameters of a job. These parameters could include how long a job should run (walltime), what resources are necessary to run, and what to execute. The following is an example submit file:

```
#PBS -N localBlast
#PBS -S /bin/sh
#PBS -l nodes=1:ppn=2,walltime=240:00:00
#PBS -M user@my.organization.com
#PBS -m ea
source ~/.bashrc
cd $HOME/work/dir
sh myBlast.sh -i -v
```

This submit script specifies the name of the job (localBlast), what environment to use (/bin/sh), that it needs both processors on a single node (nodes=1:ppn=2), that it will run for at most 10 days, and that Torque should email "user@my.organization.com" when the job exits or aborts. Additionally, the user specifies where and what to execute.

PBS Options

| | |
|-----------------------|---|
| #PBS -N myJob | Assigns a job name. The default is the name of PBS job script. |
| #PBS -l nodes=4:ppn=2 | The number of nodes and processors per node. |
| #PBS -q queueName | Assigns the queue your job will use. |
| #PBS -o mypath/my.out | The path and file name for standard output. |
| #PBS -e mypath/my.err | The path and file name for standard error. |
| #PBS -j oe | Merges the standard error and output stream of the job. |
| #PBS -m b | Sends mail to the user when the job begins. |
| #PBS -m e | Sends mail to the user when the job ends. |
| #PBS -m a | Sends mail to the user when job aborts (with an error). |
| #PBS -r n | Indicates that a job should not rerun if it fails. |
| #PBS -V | Exports all environment variables to the job. |
| #PBS -S | Specifies environment |
| #PBS -m ba | Allows a user to have more than one command with the same flag by grouping the messages together on one line, else only the last command gets executed. |

| | |
|---|--|
| <code>#PBS -l walltime=01:00:00</code> | The maximum wall-clock time during which this job can run. |
| <code>#PBS -W stagein=file_list</code> | Copies the file onto the execution host before the job starts. |
| <code>#PBS -W stageout=file_list</code> | Copies file from the execution host after the job completes. |
| <code>#PBS -M usr@domain.com</code> | Specifies e-mail address to send to |

Testing

We keep our jobs at `/craneNest/Jobs`

Go to Jobs directory: `cd /craneNest/Jobs`

Check that node status for all nodes is free:

```
pbsnodes
```

if one of the nodes has a problem SSH to it and restart the mom service: `systemctl restart pbs_mom.service`, if there is problem with SSH - fix it - Torque needs SSH to work properly.

Submit a job: `qsub hello.pbs` (or other job you have)

Check status of job: `qstat`

If a job is running it has an R under "S"

If status is Q try to forcefully run it: `qrun <job_number>`

If `qrun` works (you see R under job status using `qstat`) then there might be a problem with the scheduler. restart scheduler:

```
systemctl restart pbs_sched.service
```

When job ends look at the files in your directory, you should see two new files - `test.out` `test.err`. Make you have the correct output in `test.out` and that `test.err` doesn't have problematic errors.

Note: you can edit `hello.pbs` to specify the number of nodes and CPUs on each with PPN.

```
#PBS -l nodes=14:ppn=1,walltime=199:0:30
```

Test using different number of nodes and PPN to see that everything is OK.

NOTICE

Make sure that `/etc/hosts` `/etc/hostname` `/etc/sysconfig/network` `/var/spool/torque/server_name` `/var/spool/torque/server_name.new` **ALL** have the same hostname formats in both nodes and head node!

If running jobs has problems, check `qsub -f`

Also check if jobs are stuck in node in `/var/spool/torque/undelivered` or in `/var/spool/torque/spool`

Make sure you have permissions to write to output destinations and make sure users can SSH and SCP without a password between node and head node

LDAP

(Lightweight Directory Access Protocol)

LDAP is a user management system that allows you to create users and groups that are universal for your domain. The main system is installed on a Domain Controller and client versions are installed on each of the computers in the domain.

In our case the Domain Controller is head node and client versions are installed on each of the compute nodes.

NOTE:

In our case our domain name is `cranezone` (we set this when we set the DNS)

Installation^[1]

Install the following packages:

```
$ yum install -y openldap openldap-clients openldap-servers
migrationtools
```

Generate a **LDAP** password from a secret key and save an SSHA encryption of it in `/etc/openldap/passwd` (we also saved the non-encrypted password there)

```
$ slappasswd -s <admin_password> -n > /etc/openldap/passwd
```

Generate a X509 certificate valid for **365** days (enter all blanks except for Common Name):

```
$ openssl req -new -x509 -nodes -out /etc/openldap/certs/cert.pem \
-keyout /etc/openldap/certs/priv.pem -days 365
```

Generating a 2048 bit RSA private key

.....+++

.....+++

writing new private key to '/etc/openldap/certs/priv.pem'

You are about to be asked to enter information that will be incorporated

into your certificate request.

What you are about to enter is what is called a Distinguished Name or a DN.

There are quite a few fields but you can leave some blank
For some fields there will be a default value,
If you enter '.', the field will be left blank.

Country Name (2 letter code) [XX]:
State or Province Name (full name) []:
Locality Name (eg, city) [Default City]:
Organization Name (eg, company) [Default Company Ltd]:
Organizational Unit Name (eg, section) []:
Common Name (eg, your name or your server's hostname)
[]:<serverHostname>.<domain>
Email Address []:

Secure the content of the /etc/openldap/certs directory:

```
$ cd /etc/openldap/certs
$ chown ldap:ldap *
$ chmod 600 priv.pem
```

Prepare the LDAP database:

```
$ cp /usr/share/openldap-servers/DB_CONFIG.example
/var/lib/ldap/DB_CONFIG
```

Generate database files (don't worry about error messages!):

```
$ slaptest
53d61aab hdb_db_open: database "dc=my-domain,dc=com":
db_open(/var/lib/ldap/id2entry.bdb) failed: No such file or directory
(2).
53d61aab backend_startup_one (type=hdb, suffix="dc=my-domain,dc=com"):
bi_db_open failed! (2)
slap_startup failed (test would succeed using the -u switch)
```

Change LDAP database ownership:

```
$ chown ldap:ldap /var/lib/ldap/*
```

Activate the **slapd** service at boot:

```
$ systemctl enable slapd
```

Start the **slapd** service:

```
$ systemctl start slapd
```

Check the **LDAP** activity:

```
# netstat -lt | grep ldap
tcp        0      0
0.0.0.0:ldap        0.0.0.0:*        LISTEN
tcp6       0      0
[::]:ldap        [::]:*        LISTEN
```

To start the configuration of the **LDAP** server, add the **cosine & nis LDAP** schemas:

```
$ cd /etc/openldap/schema
$ ldapadd -Y EXTERNAL -H ldapi:/// -D "cn=config" -f cosine.ldif
SASL/EXTERNAL authentication started
SASL username: gidNumber=0+uidNumber=0,cn=peercred,cn=external,cn=auth
SASL SSF: 0
adding new entry "cn=cosine,cn=schema,cn=config"
$ ldapadd -Y EXTERNAL -H ldapi:/// -D "cn=config" -f nis.ldif
SASL/EXTERNAL authentication started
SASL username: gidNumber=0+uidNumber=0,cn=peercred,cn=external,cn=auth
SASL SSF: 0
adding new entry "cn=nis,cn=schema,cn=config"
```

Then, create the `/etc/openldap/changes.ldif` file and paste the following lines:

NOTE:

There can be as little as one `dc` and as many as needed

For domain name `wits.ac.za` you must put `dc=wits,dc=ac,dc=za`

In our example our domain is `cranezone` so we use only `dc=cranezone`

You can retrieve the SSHA password from the previously saved file `/etc/openldap/passwd`

```
dn: olcDatabase={2}hdb,cn=config
```

```
changetype: modify
```

```
replace: olcSuffix
```

```
olcSuffix: dc=cranezone
```

```
dn: olcDatabase={2}hdb,cn=config
```

```
changetype: modify
```

```
replace: olcRootDN
```

```
olcRootDN: cn=Manager,dc=cranezone
```

```
dn: olcDatabase={2}hdb,cn=config
```

```
changetype: modify
```

```
replace: olcRootPW
```

```
olcRootPW: {SSHA}l8A+0c+lRcymtWuIFbbc3EJ1PRZz9mGg
```

```
dn: cn=config
```

```
changetype: modify
```

```
replace: olcTLSCertificateFile
```

```
olcTLSCertificateFile: /etc/openldap/certs/cert.pem
```

```
dn: cn=config
```

```
changetype: modify
```

```
replace: olcTLSCertificateKeyFile
```

```
olcTLSCertificateKeyFile: /etc/openldap/certs/priv.pem
```

```
dn: cn=config
```

```
changetype: modify
```

```
replace: olcLogLevel
```

```
olcLogLevel: -1
```

```
dn: olcDatabase={1}monitor,cn=config
```

```
changetype: modify
replace: olcAccess
olcAccess: {0}to * by
dn.base="gidNumber=0+uidNumber=0,cn=peercred,cn=external,cn=auth" read
by dn.base="cn=Manager,dc=cranezone" read by * none
```

Send the new configuration to the slapd server:

```
$ ldapmodify -Y EXTERNAL -H ldapi:/// -f /etc/openldap/changes.ldif
SASL/EXTERNAL authentication started
SASL username: gidNumber=0+uidNumber=0,cn=peercred,cn=external,cn=auth
SASL SSF: 0
modifying entry "olcDatabase={2}hdb,cn=config"
modifying entry "olcDatabase={2}hdb,cn=config"
modifying entry "olcDatabase={2}hdb,cn=config"
modifying entry "cn=config"
modifying entry "cn=config"
modifying entry "cn=config"
modifying entry "olcDatabase={1}monitor,cn=config"
```

Create the /etc/openldap/base.ldif file and paste the following lines:

```
dn: dc=cranezone
dc: example
objectClass: top
objectClass: domain

dn: ou=People,dc=cranezone
ou: People
objectClass: top
objectClass: organizationalUnit

dn: ou=Group,dc=cranezone
ou: Group
objectClass: top
```



```
objectClass: organizationalUnit
Build the structure of the directory service:
$ ldapadd -x -w <admin_password> -D cn=Manager,dc=cranezone
-f /etc/openldap/base.ldif
adding new entry "dc=example,dc=com"
adding new entry "ou=People,dc=example,dc=com"
adding new entry "ou=Group,dc=example,dc=com"
```

Create two users for testing:

```
$ mkdir /home/guests
$ useradd -d /<home_directory> ldapuser01 ldapuser01
$ passwd ldapuser01
Changing password for user ldapuser01.
New password:
Retype new password:
passwd: all authentication tokens updated successfully.
$ useradd -d /home/guests/ldapuser02 ldapuser02
$ passwd ldapuser02
Changing password for user ldapuser02.
New password:
Retype new password:
passwd: all authentication tokens updated successfully.
```

User Account Migration

Go to the directory for the migration of the user accounts:

```
$ cd /usr/share/migrationtools
```

Edit the **migrate_common.ph** file and replace in the following lines:

```
$DEFAULT_MAIL_DOMAIN = "cranezone";
$DEFAULT_BASE = "dc=cranezone";
```

Create the current users in the directory service:

```
$ grep ":10[0-9][0-9]" /etc/passwd > passwd
```

```
$ ./migrate_passwd.pl passwd users.ldif
$ ldapadd -x -w <admin_password> -D cn=Manager,dc=cranezone -f
users.ldif
adding new entry "uid=ldapuser01,ou=People,dc=example,dc=com"
adding new entry "uid=ldapuser02,ou=People,dc=example,dc=com"
$ grep ":10[0-9][0-9]" /etc/group > group
$ ./migrate_group.pl group groups.ldif
$ ldapadd -x -w <admin_password> -D cn=Manager,dc=cranezone -f
groups.ldif
adding new entry "cn=ldapuser01,ou=Group,dc=example,dc=com"
adding new entry "cn=ldapuser02,ou=Group,dc=example,dc=com"
```

Test the configuration with the user called **ldapuser01**:

```
$ ldapsearch -x cn=ldapuser01 -b dc=cranezone
```

Firewall Configuration

Add a new service to the firewall (**ldap**: port **tcp 389**):

```
$ firewall-cmd --permanent --add-service=ldap
```

Reload the firewall configuration:

```
# firewall-cmd --reload
```

Edit the **/etc/rsyslog.conf** file and add the following line:

```
local4.* /var/log/ldap.log
```

Restart the **rsyslog** service:

```
$ systemctl restart rsyslog
```

LDAP Client Configuration^[2]

As the **authconfig-tui** is deprecated, to configure the LDAP client side, there are two available options: **nssld** and **sssd**.

In this tutorial, the **nssld** option will be used, see the [authconfig tutorial](#) for the **sssd** option.

Install the following packages:

```
$ yum install -y openldap-clients nss-pam-ldapd
```

Then, type:

NOTE: our server hostname is crane and our domain name is cranezone

```
$ authconfig --enableforcelegacy --update
$ authconfig --enableldap --enableldapauth
--ldapserver="crane.cranezone"
--ldapbasedn="dc=cranezone"
--update
```

NOTE 1: According to your requirements, you can specify the `--enablemkhomedir` option. This option creates a local user home directory at the first connection if none exists.

NOTE 2: Type `authconfig -help | grep ldap` to remember the necessary options.

Put the **LDAP** server certificate into the `/etc/openldap/cacerts` directory:

```
$ scp root@crane.cranezone:/etc/openldap/certs/cert.pem
/etc/openldap/cacerts/cert.pem
```

Apply the correct SELinux context to the certificate:

```
$ restorecon /etc/openldap/cacerts/cert.pem
```

Activate the TLS option:

```
$ authconfig --enableldaptls --update
```

Test the configuration:

```
$ getent passwd ldapuser02
ldapuser02:*:1001:1001:ldapuser02:/home/guests/ldapuser02:/bin/bash
```

NFS Server Configuration

To get the home directory mounted, you need to configure a NFS server. The NFS server is called `instructor.example.com` in the procedure.

NOTE: It's not required to have the LDAP server and the NFS server on the same machine, it's only easier.

Automounter Client configuration

Install the following packages:

```
$ yum install -y autofs nfs-utils
```

Create a new indirect `/etc/auto.guests` map and paste the following line:

```
* -rw,nfs4 instructor.example.com:/home/guests/&
```

Add the following line at the beginning of the `/etc/auto.master` file:

```
/home/guests /etc/auto.guests
```

Start the Automounter daemon and enable it at boot:

```
$ systemctl enable autofs && systemctl start autofs
```

Test the configuration:

```
$ su - ldapuser02
```

Create New User

```
$ cd /usr/share/migrationtools
$ useradd -d /<home_directory> ldapuser01
$ passwd ldapuser01
$ grep ":10[0-9][0-9]" /etc/passwd > passwd
$ ./migrate_passwd.pl passwd users.ldif
$ ldapadd -x -w <admin_password> -D cn=Manager,dc=cranezone -f
users.ldif
```

Create New Group

```
$ cd /usr/share/migrationtools
$ groupadd <group_name>
$ grep ":10[0-9][0-9]" /etc/group > group
```

```
$ ./migrate_group.pl group groups.ldif
$ ldapadd -x -w <admin_password> -D cn=Manager,dc=cranezone
-f groups.ldif
```

Wake On LAN

Configuration

On machines that will be woken up

```
vim /etc/sysconfig/network-scripts/ifcfg-<NIC> (eno1 in our case)
```

Add the line:

```
ETHTOOL_OPTS="wol g"
```

Get MAC address from each computer and write it down:

```
ifconfig | grep ether
```

In BIOS of each computer enable Wake-On-Lan:

- Gigabyte:
 - o press F2 on boot and enter BIOS
 - o under power Power tab Enable Wake On Lan option
- HP
 - o Press F10 on boot and enter bios
 - o Advanced >> Power-On Options >> Change "Remote Wake up Boot source" to "Remote Server"

Execution

On Machine that sends wake up signal:

```
$ ether-wake <MAC_ADDRESS>
```

Testing

Shutdown remote node

send ether-wake command from head node

see that the remote node turns on.

Format and Mounting HDD

When adding a HDD (Hard Disk Drive) to a machine you first need to format it, then you mount it.

In our case we wanted `/craneNest` to be on a separate, large, HDD. This is meant for organization and to avoid hard drive memory shortages. `/craneNest` is located on `/dev/sdb1` which is the single partition of a 1TB HDD.

Connect HDD to power and SATA on the computer.

Format

Check if the computer can see your HDD:

```
$ cat /proc/partitions
```

See if there is a new HDD (usually `sdb`) and see that its size is the same as the one you put in.

Note: `sdb#` are partitions on the HDD, we will remove them.

Remove Partitions

```
$ fdisk /dev/<new HDD>
```

If the new HDD has been previously used:

make new partition table: `o`

make new partition: `n`

choose primary partition: `p`

choose partition number: `1`

choose default settings twice: `enter`, `enter`

write the new partition table: `w`

```
mkfs.xfs /dev/sdb1
```

Mount

Create folder to be mounted to:

```
$ mkdir <folder_path_and_name>
```

```
$ mount /dev/sdb1
```

Edit `fstab` so it will be mounted on boot:

```
$ vim /etc/fstab
```

Add the following row:

```
/dev/sdb1 /craneNest xfs defaults 0 1
```

NOTE

If you are remounting an old broken mount, make sure to unmount the broken mount using the command:

```
$ umount
```


Additional Scripts

We have created some scripts for our own use, they are located in `/craneNest/admin/`.

`/craneNest/admin` is added to the `$PATH` environment variable in the root users `.bashrc` on head node and as such automatically exists in the `$PATH` variable for root.

network_config

This script sets the network table to create the internal LAN. It splits the network adapter logically so that one logical network card is connected to the normal Wits network through the wits proxy and gets an IP from the wits DHCP and the other logical card works with the internal LAN with static IP settings.

It is called in `/etc/profile` and runs automatically when a user logs in to head node.

perform

This script lets you run a command on any nodes you choose, or all the nodes, using one line. The format for it is:

```
$ ./perform "<command_between_quotations>" -on  
<node numbers with spaces as separator, head for head node, all  
for all not including head node, node1 - node2 for all lists  
between node1 and node2>
```

Example:

```
$ ./perform "ll | grep hello" -on 1 2 7 - 10 13 head
```

NOTE: This script needs to be edited when adding nodes.

ping_all

Pings all the nodes.

NOTE: This script needs to be edited when adding nodes.

wits_proxy

This script sets the wits proxy settings for wits.

It is called in `/etc/profile` and runs automatically when a user logs in.

wake_cranes

This script runs the wake-on-lan command on the nodes specified. The wake-on-lan needs to be enabled in the bios on the nodes for this script to work (all nodes have been enabled as of writing this manual).

The format for it is:

```
wake_cranes <node1> <node2> ...
```

Or

```
wake_nodes all
```

NOTE: This script needs to be edited when adding nodes.

LDAP Scripts

Inside the folder `/craneNest/admin/ldap` you will find the following scripts:

- `installClient` - this script installs and adds the necessary certificate needed for ldap to work on a new client (node). This should be run on any new node.
- `newUser <user_name>` - creates a new user in the ldap database (as well as on head node, which is the domain controller) with a user home folder located in `/craneNest/admin/userAccounts` and sets the size limit on the folder (using Quota)
- `newGroup <group_name>` - creates a new group in the LDAP database.
- `deleteUser <user_name>` - deletes a user from the ldap database (as well as on head node, which is the domain controller)
- `DeleteGroup <group_name>` - deletes a LDAP group

Ganglia^[1]^[2]

Installation

Head Node

```
$ yum install epel-release
$ yum install ganglia rrdtool ganglia-gmetad ganglia-gmond ganglia-web
```

Nodes

```
$ yum install epel-release
$ yum install ganglia-gmond
$ yum update && yum
```

Configuration

Head Node

```
$ cd /etc/ganglia
$ vim gmetad.conf
```

find the line “data source” in the file and change it to:

```
data_source "<cluster_name>" 1 <head node_hostname>
```

Nodes

```
vim /etc/ganglia/gmond.conf
```

make following sections look like and add “udp_send_channel” sections for every node (including head node):

```
cluster {
  name = "bu<cluster_name (same as in gmetad.conf)>hpc"
  owner = "unspecified"
  latlong = "unspecified"
  url = "unspecified"
}

udp_send_channel {
  host = <node_hostname>
  port = 8649
  ttl = 1
}

udp_recv_channel {
  port = 8649
  retry_bind = true
}
```

Cloning Nodes

When getting a new computer node, it is much easier to copy a HDD from an existing node onto the new nodes HDD, this is called cloning. We used Clonezilla, an open source program, to do this.

Create bootable flash drive

If you have a windows computer, download YUMI and the clonezilla iso file. Use YUMI to burn the clonezilla iso onto the flash drive.

Creating Image ^[1]

Connect the flash drive to an existing node and boot from it onto clonezilla. Connect a second flash drive to save the image on.

Use default settings when clonezilla boots.

Choose language English.

Don't touch keymap

Start Clonezilla

device-image

local_dev

choose the flash drive to save image on

choose directory name of flash drive to save in

Beginner mode

Savedisk

choose image name

choose the hdd to copy (usually sda)

skip

Yes, check saved images

Not to encrypt image

make sure everything is correct. "y" to continue "n" to abort

After everything is done, poweroff

Restoring an image ^[4]

Connect the flash drive with clonezilla on it and the flash drive with the image on it to the computer. Boot from clonezilla.

Use default settings when clonezilla boots.

Choose language English.

Don't touch keymap

Start Clonezilla

device_image

local_dev

choose flash drive with image

choose directory image is in

Beginner

NOTE: If the size of the HDD in the new node is a different size the the original node you took the image from you may need to go to expert mode (instead of beginner) and mark -icds

restoredisk
select the image
choose target partition to put image on (usually sda1)
Enter
"y" to start writing to disk, "n" to abort"
Poweroff

New stuff

Set hostname

check hostname:

```
$ hostnamectl status
```

change hostname:

```
$ hostnamectl set-hostname <new-host-name-here>
```

Networking

To enable network card:

```
$ nmtui
```

Choose your network card.

Add an X on “automatically turn on”

reboot

~~To split network card (assuming main network card is eno1):
ifconfig eno1:0 10.0.0.9 netmask 255.255.255.0 up &~~

create a file in /etc/sysconfig/network-scripts name ifcfg-eno1:0 (where eno1 is your interface) and insert in it (changing a device, hwaddr, ipaddr and gateway appropriately):

```
DEVICE="eno1:0"  
BOOTPROTO=static  
TYPE="Ethernet"  
HWADDR=38:60:77:d6:03:e6  
IPADDR=10.0.0.100  
NETMASK=255.255.255.0  
ONBOOT=yes  
GATEWAY=10.0.0.100
```

Create SSH-Key

On computer to that you want to SSH into:

```
$ ssh-keygen -t rsa
```

(enter all the way through for defaults)

```
$ ssh-copy-id <user>@<destination>
```

Change Date on Machines

date MMDDHHMMYYYY

```
$ hwclock -systohc
```

On Addition of Nodes:

Head Node

/var/named/forward.<zonenumber>

/var/named/reverse.<zonenumber>

Turn off pbs_server with `$ qterm`

→ Add node to /var/spool/torque/server_priv/nodes

→ Restart pbs_server with `$ systemctl start pbs_server.service`

Reboot

Add udp_send_channel in /etc/ganglia/gmond.conf

Nodes

/etc/hosts

/etc/hostname

/etc/sysconfig/network

/etc/sysconfig/network-scripts/ifcfg-enol:0

Add udp_send_channel in /etc/ganglia/gmond.conf

If network adapter does not appear in `$ ifconfig` you may need to change the name of the ifcfg file in /etc/sysconfig/network-scripts (in most cases from ifcfg-enol to ifcfg-enp0s25)

References

- [1] "RHEL7: Configure a LDAP directory service for user connection," [Online]. Available: <https://www.certdepot.net/rhel7-configure-ldap-directory-service-user-connection/>.
- [2] "RHEL7: Configure a system to use an existing LDAP directory service for user and group information," [Online]. Available: <https://www.certdepot.net/rhel7-configure-system-use-existing-ldap-directory-service-user-group-information/>.
- [3] "Setting Up Real-Time Monitoring with 'Ganglia' for Grids and Clusters of Linux Servers," [Online]. Available: <http://www.tecmint.com/install-configure-ganglia-monitoring-centos-linux/>.
- [4] "How to Install Ganglia on CentOS 7," [Online]. Available: <http://www.slothparadise.com/how-to-install-ganglia-on-centos-7/>.
- [5] "Save Disk Image," [Online]. Available: http://clonezilla.org/show-live-doc-content.php?topic=clonezilla-live/doc/01_Save_disk_image.
- [6] "Restore Disk Image," [Online]. Available: http://clonezilla.org/show-live-doc-content.php?topic=clonezilla-live/doc/02_Restore_disk_image.
- [7] "Google Fonts," [Online]. Available: <https://fonts.google.com/?selection.family=Arvo|Lato|Nixie+One|Source+Code+Pro>.